


```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```


```
iris = pd.read_csv('/content/IRIS.csv')
```

```
iris.head(10)
```




	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	Iris-setosa
1	4.9	3.0	1.4	0.2	Iris-setosa
2	4.7	3.2	1.3	0.2	Iris-setosa
3	4.6	3.1	1.5	0.2	Iris-setosa
4	5.0	3.6	1.4	0.2	Iris-setosa
5	5.4	3.9	1.7	0.4	Iris-setosa
6	4.6	3.4	1.4	0.3	Iris-setosa
7	5.0	3.4	1.5	0.2	Iris-setosa
8	4.4	2.9	1.4	0.2	Iris-setosa
9	4.9	3.1	1.5	0.1	Iris-setosa

```
iris.tail(10)
```




	sepal_length	sepal_width	petal_length	petal_width	species
140	6.7	3.1	5.6	2.4	Iris-virginica
141	6.9	3.1	5.1	2.3	Iris-virginica
142	5.8	2.7	5.1	1.9	Iris-virginica
143	6.8	3.2	5.9	2.3	Iris-virginica
144	6.7	3.3	5.7	2.5	Iris-virginica
145	6.7	3.0	5.2	2.3	Iris-virginica
146	6.3	2.5	5.0	1.9	Iris-virginica
147	6.5	3.0	5.2	2.0	Iris-virginica
148	6.2	3.4	5.4	2.3	Iris-virginica
149	5.9	3.0	5.1	1.8	Iris-virginica

```
iris.shape
```

 (150, 5)

```
iris.info()
```



```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 150 entries, 0 to 149
Data columns (total 5 columns):
#   Column          Non-Null Count  Dtype
---  ---          -
0   sepal_length    150 non-null   float64
1   sepal_width     150 non-null   float64
2   petal_length    150 non-null   float64
3   petal_width     150 non-null   float64
4   species         150 non-null   object
dtypes: float64(4), object(1)
memory usage: 6.0+ KB
```

```
iris.isna().sum()
```

```

0
sepal_length 0
sepal_width 0
petal_length 0
petal_width 0
species 0

```

dtype: int64

```
iris.duplicated().sum()
```

```
3
```

```
iris= iris.drop_duplicates()
```

```
iris.shape
```

```
(147, 5)
```

```
iris
```

```

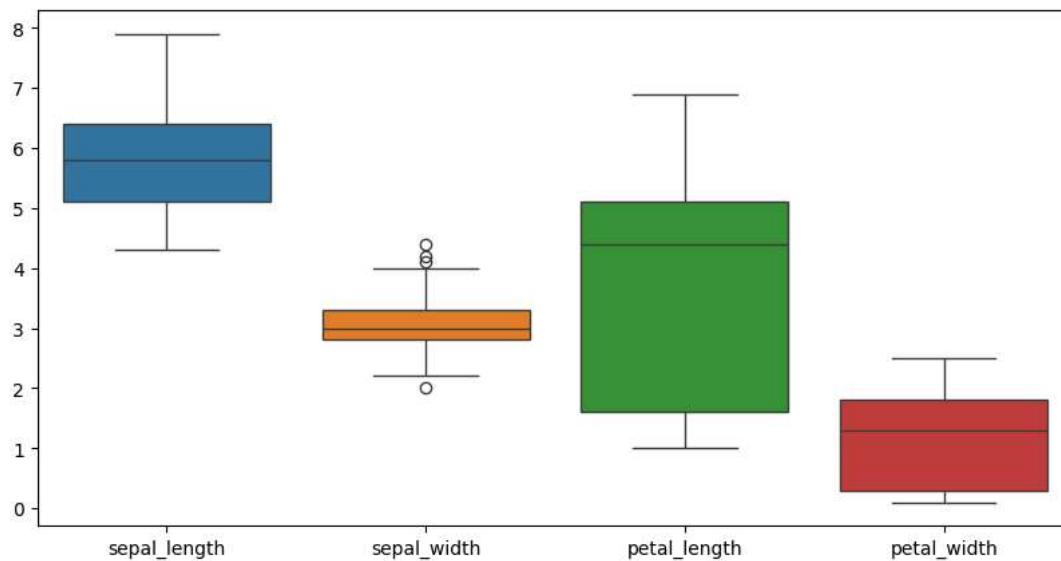
sepal_length sepal_width petal_length petal_width species
0          5.1          3.5          1.4          0.2  Iris-setosa
1          4.9          3.0          1.4          0.2  Iris-setosa
2          4.7          3.2          1.3          0.2  Iris-setosa
3          4.6          3.1          1.5          0.2  Iris-setosa
4          5.0          3.6          1.4          0.2  Iris-setosa
...         ...         ...         ...         ...         ...
145         6.7          3.0          5.2          2.3  Iris-virginica
146         6.3          2.5          5.0          1.9  Iris-virginica
147         6.5          3.0          5.2          2.0  Iris-virginica
148         6.2          3.4          5.4          2.3  Iris-virginica
149         5.9          3.0          5.1          1.8  Iris-virginica

```

147 rows × 5 columns

```
fig,ax = plt.subplots(figsize=(10,5))
sns.boxplot(iris,ax=ax)
```

```
<Axes: >
```



```
iris.describe()
```



	sepal_length	sepal_width	petal_length	petal_width
<b>count</b>	147.000000	147.000000	147.000000	147.000000
<b>mean</b>	5.856463	3.055782	3.780272	1.208844
<b>std</b>	0.829100	0.437009	1.759111	0.757874
<b>min</b>	4.300000	2.000000	1.000000	0.100000
<b>25%</b>	5.100000	2.800000	1.600000	0.300000
<b>50%</b>	5.800000	3.000000	4.400000	1.300000
<b>75%</b>	6.400000	3.300000	5.100000	1.800000
<b>max</b>	7.900000	4.400000	6.900000	2.500000

```
iris.species.unique()
```



```
array(['Iris-setosa', 'Iris-versicolor', 'Iris-virginica'], dtype=object)
```

```
iris['species'].value_counts()
```



	count
species	
<b>Iris-versicolor</b>	50
<b>Iris-virginica</b>	49
<b>Iris-setosa</b>	48

```
dtype: int64
```

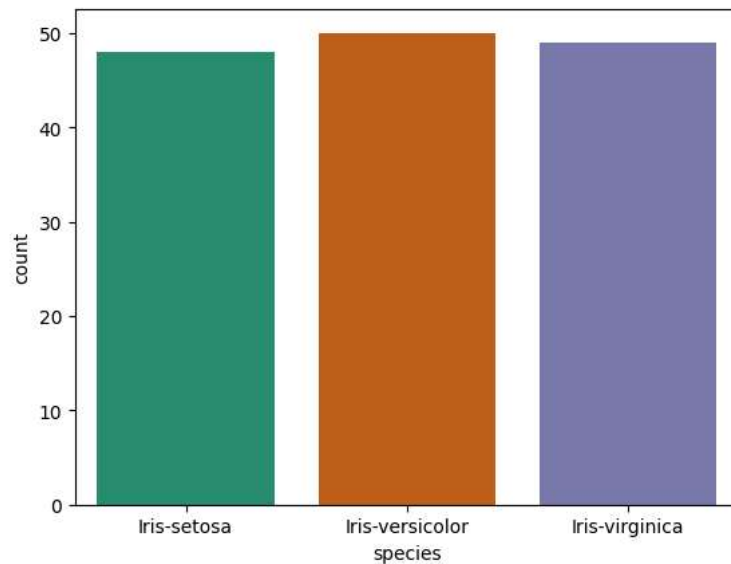
```
sns.countplot(x = "species",data=iris,palette="Dark2")
```



```
<ipython-input-18-481d8148ff53>:1: FutureWarning:
```

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend` to `False`.

```
sns.countplot(x = "species",data=iris,palette="Dark2")
<Axes: xlabel='species', ylabel='count'>
```



```
fig,ax=plt.subplots(2,2, figsize=(12,8))
```


```
ax[0,0].hist(iris['sepal_length'],bins=20,edgecolor='black')
ax[0,0].set_xlabel('sepal_length')
ax[0,0].set_ylabel('frequency')
```

```
ax[0,1].hist(iris['sepal_width'],bins=20,edgecolor='black',color='brown')
ax[0,1].set_xlabel('sepal_width')
ax[0,1].set_ylabel('frequency')
```

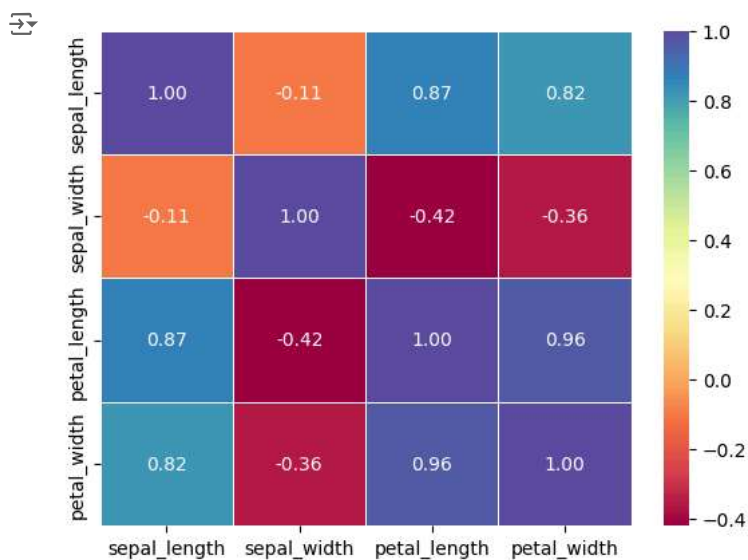
```
ax[1,0].hist(iris['petal_length'],bins=20,edgecolor='black',color='purple')
ax[1,0].set_xlabel('petal_length')
ax[1,0].set_ylabel('frequency')
```

```
ax[1,1].hist(iris['petal_width'],bins=20,edgecolor='black',color='green')
```


```
ax[1,1].set_xlabel('petal_width')
ax[1,1].set_ylabel('frequency')
plt.show()
```




```
fig,ax=plt.subplots()
sns.heatmap(iris.drop(columns='species').corr(), annot=True,cmap='Spectral',fmt=".2f",ax=ax,linewidth=0.5)
plt.show()
```




```
fig,ax=plt.subplots(figsize=(8,6))
sns.scatterplot(x="petal_width",y="petal_length",hue="species",data=iris)
plt.show()
```




```
fig,ax=plt.subplots(figsize=(8,6))
sns.scatterplot(x="sepal_width",y="sepal_length",hue="species",data=iris)
plt.show()
```

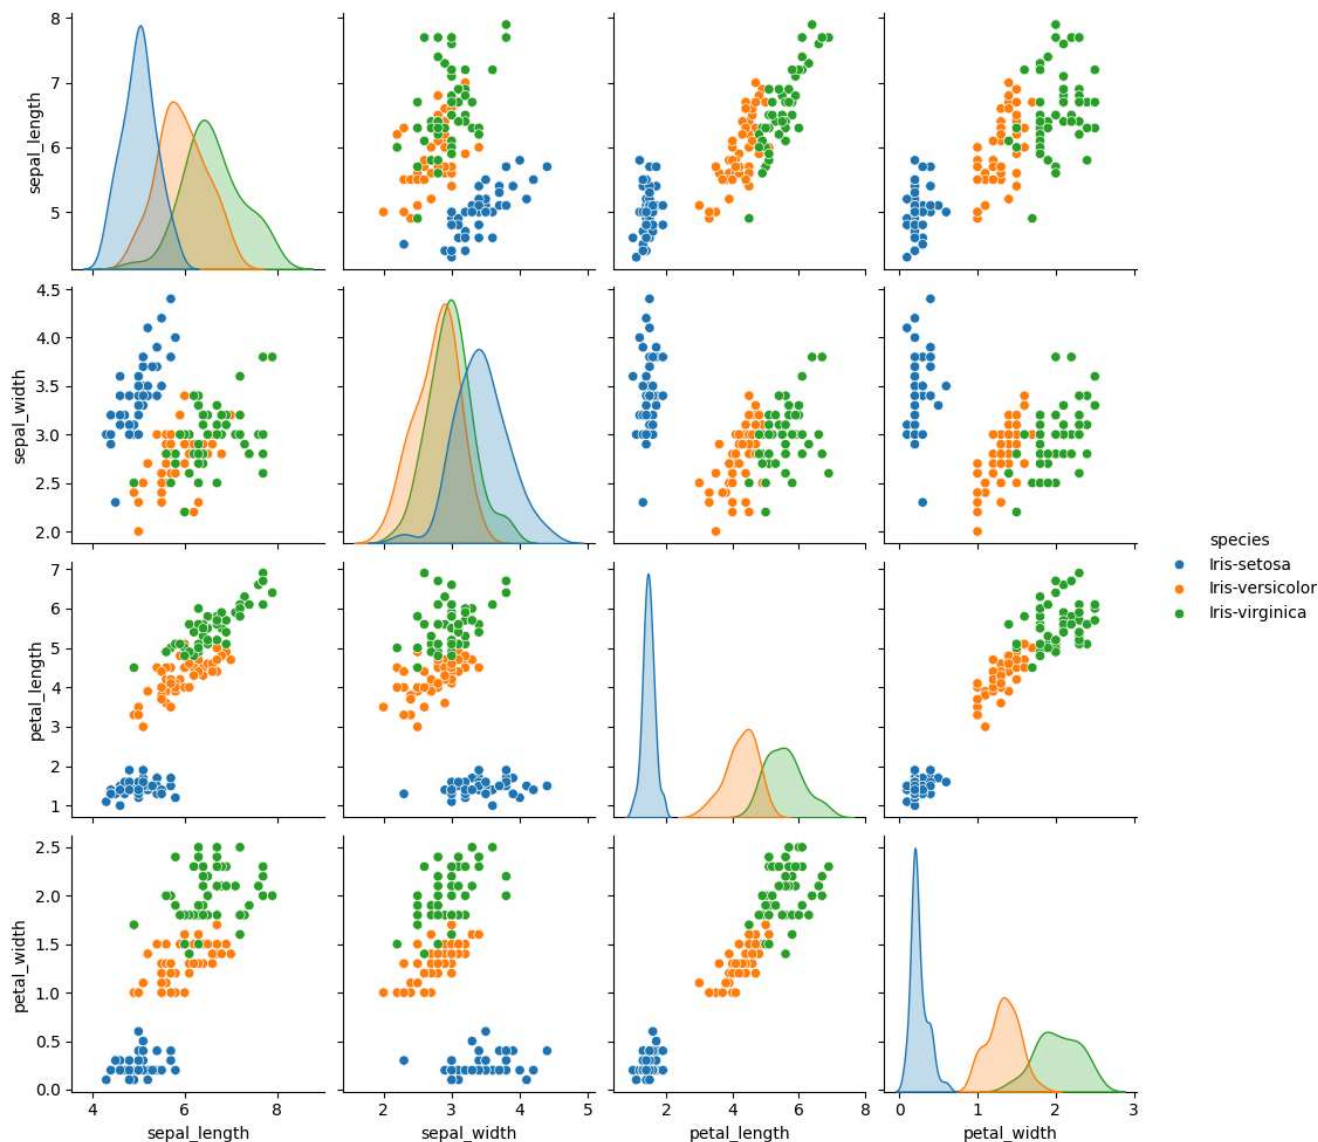


```
# Scatterplot for sepal length and petal length by species
fig,ax=plt.subplots(figsize=(8,6))
sns.scatterplot(x="petal_length",y="sepal_length",hue="species",data=iris)
plt.show()
```




```
sns.pairplot(iris, hue='species')
```

 <seaborn.axisgrid.PairGrid at 0x7be91af42b90>



```
from sklearn.preprocessing import LabelEncoder
label_encoder = LabelEncoder()
```

```
iris['species']=label_encoder.fit_transform(iris['species'])
iris
```



	sepal_length	sepal_width	petal_length	petal_width	species
0	5.1	3.5	1.4	0.2	0
1	4.9	3.0	1.4	0.2	0
2	4.7	3.2	1.3	0.2	0
3	4.6	3.1	1.5	0.2	0
4	5.0	3.6	1.4	0.2	0
...	...	...	...	...	...
145	6.7	3.0	5.2	2.3	2
146	6.3	2.5	5.0	1.9	2
147	6.5	3.0	5.2	2.0	2
148	6.2	3.4	5.4	2.3	2
149	5.9	3.0	5.1	1.8	2

147 rows × 5 columns

```

from sklearn.model_selection import train_test_split

# Define feature and target variables
x= iris.drop('species',axis=1)
y=iris['species']

#Split dataset into training and testing sets
x_train,x_test,y_train,y_test = train_test_split(x,y, test_size=0.2)

train_data= x_train.join(y_train)
train_data

```



	sepal_length	sepal_width	petal_length	petal_width	species
8	4.4	2.9	1.4	0.2	0
47	4.6	3.2	1.4	0.2	0
33	5.5	4.2	1.4	0.2	0
143	6.8	3.2	5.9	2.3	2
97	6.2	2.9	4.3	1.3	1
...	...	...	...	...	...
115	6.4	3.2	5.3	2.3	2
3	4.6	3.1	1.5	0.2	0
118	7.7	2.6	6.9	2.3	2
99	5.7	2.8	4.1	1.3	1
44	5.1	3.8	1.9	0.4	0

117 rows × 5 columns

```


from sklearn.linear_model import LogisticRegression

#Build a Logisitic Regression model
fitted_model_lr = LogisticRegression()
#Train the model
fitted_model_lr.fit(x_train,y_train)
#Make predictions
y_pred_lr = fitted_model_lr.predict(x_test)
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

# Model's performance evaluating
accuracy = accuracy_score(y_test, y_pred_lr)
print(f' Accuracy for LR: {accuracy:.4f}')
print(classification_report(y_test, y_pred_lr))
cf=confusion_matrix(y_test,y_pred_lr)
sns.heatmap(cf,annot=True,fmt=".2f",cmap="Spectral",linewidth=0.5)
plt.title('Confusion Matrix for LR Model')
plt.xlabel('Predicted labels')
plt.ylabel('True labels')
plt.show()

```



 /usr/local/lib/python3.10/dist-packages/sklearn/linear\_model/\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status= STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

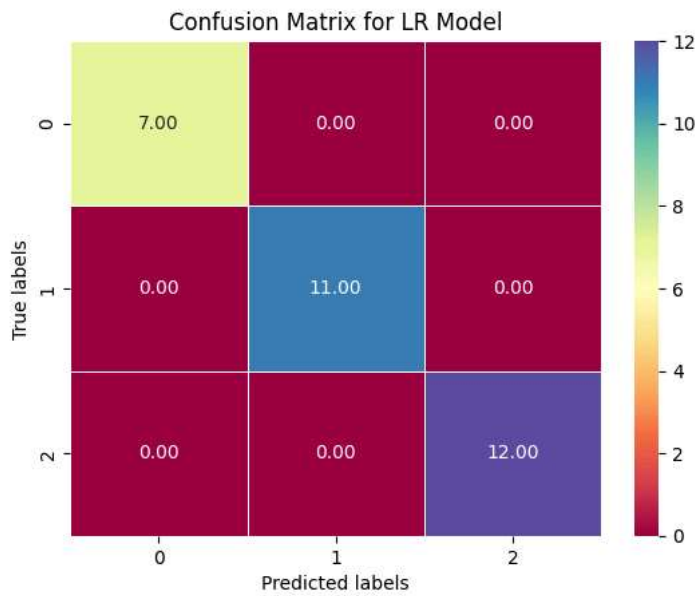
Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

n\_iter\_i = \_check\_optimize\_result(

Accuracy for LR: 1.0000

	precision	recall	f1-score	support
0	1.00	1.00	1.00	7
1	1.00	1.00	1.00	11
2	1.00	1.00	1.00	12
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

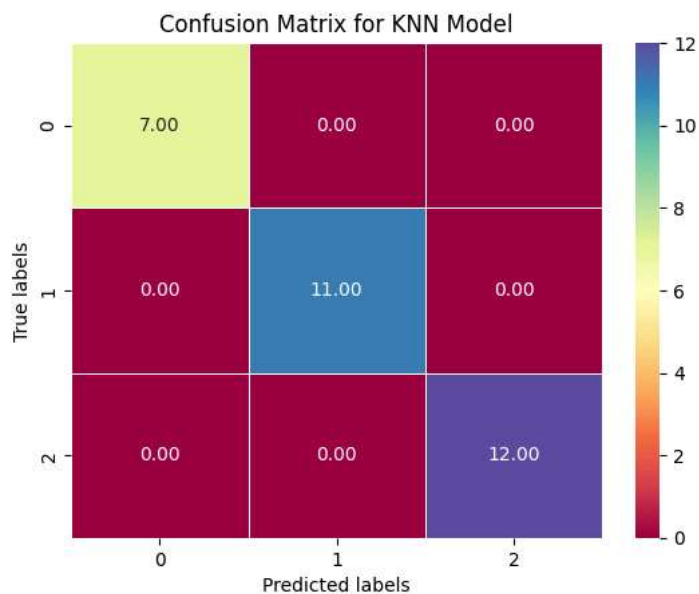


```
from sklearn.neighbors import KNeighborsClassifier
```

```
#Build a KNN model
fitted_model_knn = KNeighborsClassifier()
#Train the model
fitted_model_knn.fit(x_train,y_train)
#Make predictions
y_pred_knn = fitted_model_knn.predict(x_test)
# Model's performance evaluating
accuracy1 = accuracy_score(y_test, y_pred_knn)
print(f' Accuracy for KNN: {accuracy1:.4f}')
print(classification_report(y_test, y_pred_knn))
cf1=confusion_matrix(y_test,y_pred_knn)
sns.heatmap(cf1,annot=True,fmt=".2f",cmap="Spectral",linewidth=0.5)
plt.title('Confusion Matrix for KNN Model')
plt.xlabel('Predicted labels')
plt.ylabel('True labels')
plt.show()
```

🔗 Accuracy for KNN: 1.0000

	precision	recall	f1-score	support
0	1.00	1.00	1.00	7
1	1.00	1.00	1.00	11
2	1.00	1.00	1.00	12
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30



```
from sklearn.svm import SVC

#Build a KNN model
fitted_model_svm = SVC()
#Train the model
fitted_model_svm.fit(x_train,y_train)
#Make predictions
y_pred_svm = fitted_model_svm.predict(x_test)
# Model's performance evaluating
accuracy2 = accuracy_score(y_test, y_pred_svm)
print(f' Accuracy for SVM: {accuracy2:.4f}')
print(classification_report(y_test, y_pred_svm))
cf2=confusion_matrix(y_test,y_pred_svm)
sns.heatmap(cf2,annot=True,fmt=".2f",cmap="Spectral",linewidth=0.5)
plt.title('Confusion Matrix for SVM Model')
plt.xlabel('Predicted labels')
plt.ylabel('True labels')
plt.show()
```

```

Accuracy for SVM: 1.0000
precision    recall  f1-score   support

      0       1.00      1.00      1.00        7
      1       1.00      1.00      1.00       11
      2       1.00      1.00      1.00       12

   accuracy          1.00         30

```

from sklearn.tree import DecisionTreeClassifier

```

#Build a KNN model
fitted_model_dt = DecisionTreeClassifier()
#Train the model
fitted_model_dt.fit(x_train,y_train)
#Make predictions
y_pred_dt = fitted_model_dt.predict(x_test)
# Model's performance evaluating
accuracy3 = accuracy_score(y_test, y_pred_dt)
print(f' Accuracy for DT: {accuracy3:.4f}')
print(classification_report(y_test, y_pred_dt))
cf3=confusion_matrix(y_test,y_pred_dt)
sns.heatmap(cf3,annot=True,fmt=".2f",cmap="Spectral",linewidth=0.5)
plt.title('Confusion Matrix for DT Model')
plt.xlabel('Predicted labels')
plt.ylabel('True labels')
plt.show()

```

```

Accuracy for DT: 1.0000
precision    recall  f1-score   support

      0       1.00      1.00      1.00        7
      1       1.00      1.00      1.00       11
      2       1.00      1.00      1.00       12

   accuracy          1.00         30
  macro avg       1.00      1.00      1.00         30
 weighted avg       1.00      1.00      1.00         30

```

