


```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```


```
advertising = pd.DataFrame(pd.read_csv("advertising.csv"))
```

```
advertising.head()
```




	TV	Radio	Newspaper	Sales
0	230.1	37.8	69.2	22.1
1	44.5	39.3	45.1	10.4
2	17.2	45.9	69.3	12.0
3	151.5	41.3	58.5	16.5
4	180.8	10.8	58.4	17.9

```
advertising.shape
```

 (200, 4)

```
advertising.info()
```




```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 4 columns):
#   Column      Non-Null Count  Dtype
---  -
0    TV          200 non-null    float64
1    Radio        200 non-null    float64
2    Newspaper    200 non-null    float64
3    Sales        200 non-null    float64
dtypes: float64(4)
memory usage: 6.4 KB
```

```
advertising.describe()
```



	TV	Radio	Newspaper	Sales
count	200.000000	200.000000	200.000000	200.000000
mean	147.042500	23.264000	30.554000	15.130500
std	85.854236	14.846809	21.778621	5.283892
min	0.700000	0.000000	0.300000	1.600000
25%	74.375000	9.975000	12.750000	11.000000
50%	149.750000	22.900000	25.750000	16.000000
75%	218.825000	36.525000	45.100000	19.050000
max	296.400000	49.600000	114.000000	27.000000


```
advertising.isnull().sum()*100/advertising.shape[0]
```




	0
TV	0.0
Radio	0.0
Newspaper	0.0
Sales	0.0

dtypes: float64


```
fig, axs = plt.subplots(3, figsize = (5,5))
plt1 = sns.boxplot(advertising['TV'], ax = axs[0])
plt2 = sns.boxplot(advertising['Newspaper'], ax = axs[1])
plt3 = sns.boxplot(advertising['Radio'], ax = axs[2])
plt.tight_layout()
```




```
sns.boxplot(advertising['Sales'])  
plt.show()
```



```
sns.pairplot(advertising, x_vars=['TV', 'Newspaper', 'Radio'], y_vars='Sales', height=4, aspect=1, kind='scatter')  
plt.show()
```



```
sns.heatmap(advertising.corr(), cmap="YlGnBu", annot = True)  
plt.show()
```



```
X = advertising['TV']
y = advertising['Sales']
```

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, train_size = 0.7, test_size = 0.3, random_state = 100)
```

```
X_train.head()
```

```

   TV
74  213.4
3   151.5
185 205.0
26  142.9
90  134.3
```

```
dtype: float64
```

```
y_train.head()
```

```

   Sales
74    17.0
3     16.5
185    22.6
26    15.0
90    14.0
```

```
dtype: float64
```

```
import statsmodels.api as sm
```

```

#Add a constant to get an intercept
X_train_sm = sm.add_constant(X_train)

# Fit the regression line using 'OLS'
lr = sm.OLS(y_train, X_train_sm).fit()
```

```
lr.params
```

```

      0
const 6.948683
TV    0.054546
```

```
dtype: float64
```

```
print(lr.summary())
```



OLS Regression Results

```
=====
Dep. Variable:          Sales    R-squared:                0.816
Model:                  OLS      Adj. R-squared:           0.814
Method:                 Least Squares    F-statistic:           611.2
Date:                   Thu, 12 Sep 2024    Prob (F-statistic):    1.52e-52
Time:                   17:55:49    Log-Likelihood:        -321.12
No. Observations:       140      AIC:                   646.2
Df Residuals:           138      BIC:                   652.1
Df Model:                1
Covariance Type:        nonrobust
=====
```


	coef	std err	t	P> t	[0.025	0.975]
const	6.9487	0.385	18.068	0.000	6.188	7.709
TV	0.0545	0.002	24.722	0.000	0.050	0.059

```
=====
Omnibus:                0.027    Durbin-Watson:           2.196
Prob(Omnibus):           0.987    Jarque-Bera (JB):         0.150
Skew:                   -0.006    Prob(JB):                 0.928
Kurtosis:               2.840    Cond. No.                 328.
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

```
plt.scatter(X_train, y_train)
plt.plot(X_train, 6.948 + 0.054*X_train, 'r')
plt.show()
```



```
y_train_pred = lr.predict(X_train_sm)
res = (y_train - y_train_pred)
```

```
fig = plt.figure()
sns.distplot(res, bins = 15)
fig.suptitle('Error Terms', fontsize = 15)          # Plot heading
plt.xlabel('y_train - y_train_pred', fontsize = 15)  # X-label
plt.show()
```


```
<ipython-input-24-723b49e70e34>:2: UserWarning:
`distplot` is a deprecated function and will be removed in seaborn v0.14.0.
```

Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).


For a guide to updating your code to use the new functions, please see <https://gist.github.com/mwaskom/de44147ed2974457ad6372750bbe5751>

```
sns.distplot(res, bins = 15)
```

Error Terms



```
plt.scatter(X_train,res)
plt.show()
```



```
# Add a constant to X_test
X_test_sm = sm.add_constant(X_test)

# Predict the y values corresponding to X_test_sm
y_pred = lr.predict(X_test_sm)

y_pred.head()
```



0

126	7.374140
104	19.941482
99	14.323269
92	18.823294
111	20.132392

dtype: float64

```
from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score
```

```
np.sqrt(mean_squared_error(y_test, y_pred))
```



2.019296008966232

```
r_squared = r2_score(y_test, y_pred)
r_squared
```



0.792103160124566

```
plt.scatter(X_test, y_test)
plt.plot(X_test, 6.948 + 0.054 * X_test, 'r')
plt.show()
```

